

ASSIGNMENT-3

GAURANG GARG [102303134]

SEQUENTIAL :

Threads	Time (s)	Speedup	Efficiency
1	2.901	1.00	100%
2	2.942	0.99	49.32%
4	2.929	0.99	24.76%
8	2.884	1.01	12.58%
16	2.883	1.01	6.29%


Verification (diagonal ≈ 1.0):
result[0][0] = 1.000000 [OK]
result[1][1] = 1.000000 [OK]
result[2][2] = 1.000000 [OK]
result[3][3] = 1.000000 [OK]
result[4][4] = 1.000000 [OK]

Performance counter stats for './correlate_seq 2000 2000':

17,542.74 msec task-clock	# 1.000 CPUs utilized
197 context-switches	# 11.230 /sec
80 cpu-migrations	# 4.560 /sec
23,578 page-faults	# 1.344 K/sec
77,902,692,575 cycles	# 4.441 GHz (35.70%)
891,437,915 stalled-cycles-frontend	# 1.14% frontend cycles idle (35.71%)
109,446,233,956 instructions	# 1.40 insn per cycle
	# 0.01 stalled cycles per insn (35.71%)
12,228,277,687 branches	# 697.056 M/sec (35.72%)
22,985,444 branch-misses	# 0.19% of all branches (35.73%)
24,317,715,568 L1-dcache-loads	# 1.386 G/sec (35.73%)
3,315,035,428 L1-dcache-load-misses	# 13.63% of all L1-dcache accesses (35.73%)
<not supported> LLC-loads	
<not supported> LLC-load-misses	
226,149,852 L1-icache-loads	# 12.891 M/sec (35.73%)
227,093 L1-icache-load-misses	# 0.10% of all L1-icache accesses (35.72%)
65,742,868 dTLB-loads	# 3.748 M/sec (35.71%)
57,547,718 dTLB-load-misses	# 87.53% of all dTLB cache accesses (35.71%)
17,245 iTLB-loads	# 983.028 /sec (35.70%)
78,388 iTLB-load-misses	# 454.55% of all iTLB cache accesses (35.70%)
3,185,149,568 L1-dcache-prefetches	# 181.565 M/sec (35.70%)
<not supported> L1-dcache-prefetch-misses	

17.549866266 seconds time elapsed

17.480921000 seconds user
0.062473000 seconds sys



OPENMP :

Threads	Time (s)	Speedup	Efficiency
1	2.786	1.00	100%
2	1.419	1.96	98.20%
4	0.736	3.79	94.66%
8	0.377	7.39	92.40%
10	0.309	9.02	90.17%
12	0.262	10.64	88.65%
14	0.231	12.04	86.01%
16	0.220	12.64	79.03%

```

Verification (diagonal ≈ 1.0):
result[0][0] = 1.000000 [OK]
result[1][1] = 1.000000 [OK]
result[2][2] = 1.000000 [OK]
result[3][3] = 1.000000 [OK]
result[4][4] = 1.000000 [OK]


Performance counter stats for './correlate_omp 2000 2000':

```

Event	Value	Description
task-clock	28,656.89 msec	4.268 CPUs utilized
context-switches	1,562	# 54.507 /sec
cpu-migrations	90	# 3.141 /sec
page-faults	23,646	# 825.142 /sec
cycles	116,600,033,531	# 4.069 GHz (35.74%)
stalled-cycles-frontend	906,035,737	# 0.78% frontend cycles idle (35.77%)
instructions	164,054,949,816	# 1.41 insn per cycle
		# 0.01 stalled cycles per insn (35.76%)
branches	18,425,957,522	# 642.985 M/sec (35.73%)
branch-misses	31,349,447	# 0.17% of all branches (35.72%)
L1-dcache-loads	36,585,892,284	# 1.277 G/sec (35.74%)
L1-dcache-load-misses	6,659,529,260	# 18.20% of all L1-dcache accesses (35.73%)
LLC-loads	<not supported>	
LLC-load-misses	252,278,041	L1-icache-loads # 8.803 M/sec (35.71%)
	410,149	L1-icache-load-misses # 0.16% of all L1-icache accesses (35.69%)
dTLB-loads	110,907,253	# 3.870 M/sec (35.68%)
dTLB-load-misses	87,556,795	# 78.95% of all dTLB cache accesses (35.65%)
iTLB-loads	43,335	# 1.512 K/sec (35.65%)
iTLB-load-misses	81,982	# 189.18% of all iTLB cache accesses (35.70%)
L1-dcache-prefetches	6,320,194,312	# 220.547 M/sec (35.73%)
L1-dcache-prefetch-misses	<not supported>	

6.714600069 seconds time elapsed

28.538695000 seconds user
0.112789000 seconds sys



OPTIMIZED :

Threads	Time (s)	Speedup	Efficiency
1	1.236	1.00	100%
2	0.633	1.95	97.58%
4	0.331	3.73	93.28%
8	0.189	6.55	81.82%
10	0.150	8.27	82.66%
12	0.126	9.84	81.97%
14	0.112	11.01	78.62%
16	0.124	10.01	62.53%


```
Verification (diagonal ~ 1.0):
result[0][0] = 1.000000 [OK]
result[1][1] = 1.000000 [OK]
result[2][2] = 1.000000 [OK]
result[3][3] = 1.000000 [OK]
result[4][4] = 1.000000 [OK]
```

Performance counter stats for './correlate_opt 2000 2000':

14,263.33 msec task-clock	# 4.570 CPUs utilized
661 context-switches	# 46.343 /sec
48 cpu-migrations	# 3.365 /sec
23,650 page-faults	# 1.658 K/sec
56,434,384,475 cycles	# 3.957 GHz (35.74%)
1,825,892,577 stalled-cycles-frontend	# 3.24% frontend cycles idle (35.74%)
46,579,578,883 instructions	# 0.83 insn per cycle
	# 0.04 stalled cycles per insn (35.72%)
9,276,462,090 branches	# 650.371 M/sec (35.75%)
25,976,751 branch-misses	# 0.28% of all branches (35.76%)
23,504,430,868 L1-dcache-loads	# 1.648 G/sec (35.74%)
6,507,629,356 L1-dcache-load-misses	# 27.69% of all L1-dcache accesses (35.75%)
<not supported> LLC-loads	
<not supported> LLC-load-misses	
231,682,368 L1-icache-loads	# 16.243 M/sec (35.74%)
403,872 L1-icache-load-misses	# 0.17% of all L1-icache accesses (35.68%)
110,012,378 dTLB-loads	# 7.713 M/sec (35.65%)
86,524,415 dTLB-load-misses	# 78.65% of all dTLB cache accesses (35.66%)
25,841 iTLB-loads	# 1.812 K/sec (35.67%)
85,682 iTLB-load-misses	# 331.57% of all iTLB cache accesses (35.67%)
4,723,325,029 L1-dcache-prefetches	# 331.152 M/sec (35.71%)
<not supported> L1-dcache-prefetch-misses	

3.120876142 seconds time elapsed

14.166922000 seconds user
0.094999000 seconds sys



OBSERVATION :

Version 1 - Sequential Baseline :

The baseline took **~2.95 seconds** with 1 thread. When you gave it more threads (2, 4, 8... 16), the time barely changed it stayed around 2.8-3.1 seconds. That's because Version 1 has no parallel code. Adding more threads did nothing; the work still ran on one core. All "speedups" are basically 1x (no improvement at all).

Bottom line: More threads = no benefit. Works correctly, but uses only 1 core.

Version 2 - OpenMP Parallel :

Baseline here was **~2.79 seconds** with 1 thread. With more threads, it got noticeably faster:

Threads	Time	Speedup
1	2.79s	1x
2	1.42s	~2x
4	0.74s	~4x
8	0.38s	~7.4x
16	0.22s	~12.6x

This scales really well. Doubling the threads roughly halves the time. The efficiency (how well each thread is being used) stays above 85-90% until 14 threads, then drops slightly at 16. That's excellent parallel scaling.

Bottom line: Adding threads genuinely helps. Near perfect scaling up to ~12 threads.

Version 3 - Optimized (AVX2 SIMD + OpenMP) :

Baseline here was just **~1.24 seconds** that's already 2.25x faster than Version 2 with just 1 thread. This is because AVX2 does 4 floating-point operations in one CPU instruction instead of 1.

With threads:

Threads	Time	Speedup
1	1.24s	1x
8	0.19s	~6.5x
12	0.13s	~9.8x
16	0.12s	~10x

Notice at 16 threads the speedup actually drops a little compared to 14 threads (from 11x to 10x). This is called diminishing returns the overhead of managing 16 threads starts eating into the gains.

Bottom line: Fastest version overall. But above 12-14 threads, adding more threads stops helping.

The perf stat Numbers :

Metric	Seq	OMP	Optimized
Total CPU time	~26.6s	~28.7s	~14.3s
Instructions per cycle	1.39	1.41	0.83
L1 cache miss rate	13.6%	18.2%	27.7%

Instructions per cycle (IPC): Version 3 has a lower IPC (0.83 vs 1.39), which sounds bad, but it's misleading. Each AVX2 instruction does 4x the math. So fewer instructions doing more work is actually better.

Cache miss rate: Version 3 misses L1 cache more (27.7%) because AVX2 pulls in larger data chunks at a time, sometimes exceeding what fits in the fast L1 cache. This is a known tradeoff you gain from SIMD but pay a little in cache pressure.

dTLB misses (~87%): All three versions have very high TLB miss rates. This means the program accesses memory spread across many different pages, which is expected for a large 2000×2000 matrix. This is a minor bottleneck you can't easily avoid.